

Operating Systems

Course Code: **71203002004**

***Kernel I/O Subsystem – Services & Implementation***

*by -*

*Minal Rajwar*



# Caching, Spooling, and Device Reservation

## Spooling (Simultaneous Peripheral Operation On-Line)

- **Meaning:** Temporarily stores I/O data (like print jobs) in secondary storage before sending it to slow devices.
- **Purpose:** Keeps the CPU busy by not waiting for slow I/O devices.
- **Example:** Multiple documents sent to a printer are stored in a spool queue and printed one by one (FIFO).

# Caching, Spooling, and Device Reservation

## Why Spooling is Needed

- **Avoid CPU idle time** – CPU keeps processing other tasks while I/O happens in background.
- **Synchronize I/O** – Manages multiple print jobs without mixing output.
- **Manage speed mismatch** – Handles slow I/O devices compared to fast CPUs.
- **Enable parallelism** – Allows CPU and I/O to work simultaneously.
- **Provide reliable output** – Maintains job order and consistency.

# Caching, Spooling, and Device Reservation

## How Spooling Works

1. Data sent by CPU → stored in spool area (on disk).
2. When device is ready → data fetched → processed.
3. Uses FIFO queue for fairness and order.
4. Combines buffering (temporary storage) and queuing (job order).



# Spooling vs Buffering

Feature	Spooling	Buffering
Works on	Multiple jobs	Same job
Storage	Secondary (disk)	Main memory
Efficiency	Higher	Lower
Remote processing	Supported	Not supported
Error handling	Easier recovery (data on disk)	Buffer overflow possible
Complexity	More complex	Simpler



# Spooling

## Pros

- Improves CPU utilization and efficiency
- Manages multiple user jobs
- Ensures ordered and error-free output

## Cons

- Needs large storage
- Increases disk traffic
- Delays if spool is full

# Error Handling and I/O Protection

## Error Handling

1. **Detecting errors:** OS checks for issues (device failure, missing file).
2. **Reporting errors:** Notifies system or app via error messages or exceptions.
3. **Responding to errors:** Retries operation, uses fallback data, or safely terminates faulty process.
4. **Security:** Avoids revealing sensitive info in error messages.

# Error Handling and I/O Protection

## I/O Protection

- **Privileged Instructions:** User programs **cannot directly access hardware** — only OS can.
- **Kernel Mediation:** Kernel handles all I/O system calls.
- **Goal:** Prevent user programs from corrupting data or crashing devices.
- **Benefits:** Security, stability, and controlled access to hardware.

# Performance Optimization in Kernel I/O

Optimizing how data moves between CPU and storage.

## 1. I/O Schedulers

- **Purpose:** Decide the order of disk I/O requests for best performance.
- **Common Schedulers:**
  - **Noop:** FIFO queue, best for SSDs.
  - **Deadline:** Prioritizes reads with time limits (databases).
  - **CFQ:** Fair I/O sharing for all processes.
  - **mq-deadline:** Scalable version for modern NVMe SSDs.
- **Tuning Parameters:** Adjust values like `read_expire` or `fifo_batch`.

## 2. Asynchronous I/O (AIO) and io\_uring

- **AIO:** App continues while I/O runs in background.
- **io\_uring:** Modern Linux interface — faster, less kernel overhead.

# Performance Optimization in Kernel I/O

## 3. Caching and Buffer Management

- **Page cache & buffer cache** store frequently accessed data in RAM.
- **Tunable Parameters:**
  - `vm.dirty_ratio` – how much RAM to use before writing to disk
  - `vm.swappiness` – how aggressively to use swap

## 4. Disk and Filesystem Optimization

- **Choose right filesystem:** `ext4` (general), `XFS` (large files).
- **Use mount options:** `noatime` (skip access-time updates).
- **Proper disk alignment:** Reduces I/O delays.

# Performance Optimization in Kernel I/O

## 5. Hardware Optimization

- Use **NVMe SSDs** for speed.
- Use **RAID** for redundancy and performance.

## 6. Monitoring Tools

- **iostat, iotop, atop** – identify I/O bottlenecks and tune accordingly.

# I/O Scheduling Overview within Kernel

## What is I/O Scheduling?

- Manages **order of I/O requests** to improve performance and fairness.
- Needed because **I/O devices are slower** than CPU.

## Why Needed

- Reduce wait time
- Prevent long delays
- Optimize disk arm movement
- Fairly share device access among processes

# I/O Scheduling Overview within Kernel

## How It Works

1. **Request Initiation** – Process sends I/O request.
2. **I/O Traffic Controller** – Tracks all devices and requests.
3. **I/O Scheduler** – Decides which request runs next.
4. **Device Handler** – Transfers data.
5. **Completion** – Notifies program.

## DISCUSSION & REVISION

1. Spooling temporarily stores data in which type of memory before sending it to a device?
2. Which I/O scheduler is best suited for SSDs or virtualized environments?
3. User programs must make system calls for I/O because direct hardware access requires which type of instructions?
4. Which I/O scheduling algorithm serves requests strictly in arrival order?
5. The kernel uses which cache to store frequently accessed disk data in memory?



## REFERENCES

1. <https://www.geeksforgeeks.org/operating-systems/kernel-i-o-subsystem-in-operating-system/>
2. <https://www.tutorialspoint.com/kernel-i-o-subsystem-in-operating-system>